

Learned Latent Curves and the Hyperspherical-Harmonic Variant

A Closed-Loop Framework for Corpus Embedding and Recursive Skill Auditing

Shant Tchatalbachian

August 5, 2026

Abstract. We describe a closed-loop framework for auditing and prioritizing improvement work across a corpus of engineering artifacts (software skills, self-documentation files), and we document a sphere-aware extension of its curve-fitting stage. The framework has three composed techniques: (1) a learned latent curve, which re-expands a single 1-D ordering coordinate t into a fixed-width D -dimensional embedding via a Fourier basis with learned frequencies; (2) curve-guided recursive self-improvement, which uses sparse cells of the fitted curve as a prioritization lens for gap-mapping and bounded fixpoint-style editing cycles; and (3) a self-documentation variant of the same pipeline, retargeted at heterogeneous memory-file corpora. We give the governing equations for each stage, the empirical validation obtained on a corpus of software skills ($N = 49$ to 70), and we introduce a fourth technique — the hyperspherical-harmonic curve — which replaces the flat $[0,1]^2$ parameter manifold of (1) with the Riemann sphere S^2 (the 2-sphere, with a generalization sketch to S^N). The variant uses the canonical orthonormal hyperspherical-harmonic basis and adds a learned Möbius ϕ_θ $\text{PSL}(2, \mathbb{C})$ reparameterization of the domain. We evaluate the variant against a capacity-matched flat Fourier baseline on two splits of the corpus, report absolute holdout R^2 and the matched-parameter delta, and identify five open items that bound the claim.

1 Introduction

Problem statement.

Given N items in a K -dimensional feature space, find (i) a scalar ordering coordinate $t_i \in [0,1]$ for each item and (ii) a fixed-width D -dimensional embedding $z_i \in \mathbb{R}^D$ such that interpolation between adjacent items in the ordering is well defined and sparse regions of the ordering are detectable. This arises in corpus-audit tasks: a practitioner needs to know which parts of an ordering are thin and need more work.

This paper.

The paper has two halves. Sections 2–4 give the family background — the flat learned-latent-curve, the curve-guided recursive self-improvement pipeline that consumes it, and the self-documentation offshoot — at the level of governing equations and empirical context. Section 5 introduces the new contribution: a hyperspherical-harmonic curve on the Riemann sphere S^2 with a learned Möbius reparameterization, the sphere-aware analogue of the flat curve. Sections 6–9 give the dataset, evaluation protocol, results, and discussion for the variant. We do not claim the variant is a strict improvement over the flat curve in absolute terms; we report the absolute holdout R^2 alongside the matched-parameter delta and identify where the result holds and where it does not.

2 Background: The Learned-Latent-Curve Family

2.1 Flat curve model

For output dimension $j = 1, \dots, D$ (canonically $D = 384$) and 1-D coordinate $t \in [0,1]$, the model is

$$z_j(t) = a_{j0} + \sum_{m=1}^k (a_{j,m} \sin(2\pi f_m t) + b_{j,m} \cos(2\pi f_m t)) \quad (1)$$

with k shared learned frequencies $f_1, \dots, f_k$ and per-output coefficients $a_{j,m}, b_{j,m}$. Stacked over j , this is a curve $\gamma: \mathbb{R} \rightarrow \mathbb{R}^D$. Writing the design vector

$$\phi(t) = [1, \sin(2\pi f_1 t), \cos(2\pi f_1 t), \dots, \sin(2\pi f_k t), \cos(2\pi f_k t)] \in \mathbb{R}^{1+2k} \quad (2)$$

the model is $y(t) = C \phi(t)$ with coefficient matrix $C \in \mathbb{R}^{D \times (1+2k)}$, giving a parameter count of

$$P = k + D(1+2k) \quad (3)$$

At $D = 384$, $k = 8$: $P = 8 + 384 \times 17 = 6,536$ parameters.

2.2 Closed-form coefficient solve

With the frequencies f held fixed, the linear coefficient solve is the Tikhonov ridge

$$C = (\Phi\Phi + \lambda I)^{-1} \Phi Z \quad (4)$$

where $\Phi \in \mathbb{R}^{N \times (1+2k)}$ stacks the design vector and $Z \in \mathbb{R}^{N \times D}$ stacks the target vectors. This is the sanity floor for any gradient-descent fit at the same frequencies: a fit worse than this at the same f means the optimizer failed, not the model.

2.3 Frequency parameterization

Frequencies are stored in an unconstrained form and mapped through a softplus to guarantee strict positivity and finite gradients near zero:

$$f_m = \text{softplus}(f_m) = \log(1 + \exp(f_m)), \quad f_m > 0 \quad (5)$$

2.4 Losses

Let $t \in \mathbb{R}^N$, $Z \in \mathbb{R}^{N \times D}$ the targets, and $\hat{Z} = y(t)$ the fitted curve evaluated at the training coordinates. The total objective is

$$= \text{rec} + \text{freq} + \text{smooth} \quad (6)$$

The reconstruction term $\text{rec} = \text{mean}((\hat{Z} - Z)^2)$ measures fit. The frequency-magnitude regulariser $\text{freq} = \lambda_f \cdot \text{mean}(f^2)$ with $\lambda_f \cdot 10^{-4}$ keeps the learned frequencies from collapsing to the Nyquist limit. The smoothness regulariser

$$\text{smooth} = \lambda_s \cdot \text{mean}((y(t_{\text{grid}})[2:] - 2y(t_{\text{grid}})[1:-1] + y(t_{\text{grid}})[:-2])^2) \quad (7)$$

with t_{grid} a dense grid of ~ 512 points spanning the coordinate range and $\lambda_s \cdot 10^{-3}$, penalises the second-difference of the curve between observed points, where reconstruction loss alone is blind.

2.5 Obtaining the coordinate t

The default pipeline sets t from the top principal component of a z-scored $N \times D_{\text{feat}}$ feature matrix, mapped to $[0,1]$; a rank-uniformized variant replaces raw PC1 scores by their ranks to improve design-matrix conditioning. The explained-variance ratio of PC1 (or PC1+PC2 for a 2-D surface variant) is recorded as an explicit go/no-go gate, with $\text{PC1} \geq 0.40$ treated as evidence of usable low-rank structure.

3 Curve-Guided Recursive Self-Improvement (overview — see Appendix A for full pipeline)

3.1 Pipeline

The pipeline has four stages plus a re-fit verification step:

1. 1. Fit the curve (Stage 1) — reuse Section 2.
2. 2. Detect sparse cells (Stage 2) — a uniform 0.05×0.05 grid in $[0,1]^2$ partitions the parameter space; cells with ≤ 1 item are isolated.
3. 3. Dispatch gap-mapping (Stage 3) — a fresh-context gap-mapping subagent (focused on a single sparse-cell item, never the whole corpus) is dispatched per isolated item.
4. 4. Apply recursive-self-improvement cycles (Stage 4) — capped at 3 cycles per run, with the fixpoint rule no new substantive gaps AND old gaps closed AND no new anti-patterns.

5. 5. Re-fit and verify (Stage 5) — Stage 1 is re-run after all cycles; the closed-loop claim is verified by comparing the sparse-cell count before and after.

The success metric is

$$\Delta_{\text{sparse}} = |\text{sparse_cells}|_{\text{post}} - |\text{sparse_cells}|_{\text{pre}} < 0 \quad (8)$$

3.2 Concrete cell definition

For a 2-D surface fit, coordinates are reduced to $(u, v) \in [0,1]^2$ via PC1+PC2 of a binary primitive-coverage matrix $C \in \{0,1\}^{N \times 9}$ lifted to $D = 384$ by a fixed seeded orthonormal projection Q :

$$Z = C Q, \quad (u_i, v_i) = \text{PCA}_{1,2}(Z)_i \quad (9)$$

The $[0,1]^2$ plane is discretized into a 0.05×0.05 grid ($21 \times 21 = 441$ cells). For a cell $c = (u, v)$,

$$\text{neighbors}(c) := \{ i \text{ corpus} : \|(u_i, v_i) - c\| \leq r \}, \quad (10)$$

$$\text{is_sparse}(c) := |\text{neighbors}(c)| = 0 \quad (11)$$

with default radius $r = 0.05$.

4 Self-Documentation Variant (overview — see Appendix B for full variant)

The self-documentation offshoot retargets Eqs. 9–11 at heterogeneous memory-file corpora (SELF.md, SELF-CHANGELOG.md, and additional preference / rule files), under the granularity rule that each version is one corpus item (a changelog entry, a memory-file section, or a named row). Each corpus receives its own 9-D binary primitive basis rather than sharing one basis across structurally distinct document types, since a single shared basis was found to flatten the per-corpus signal. Near-constant columns (coverage > 0.90 or < 0.10) are dropped before the PCA lift in Eq. 9.

5 Notation

D — embedding dimension (default 384).

N — corpus size; $N_{\text{train}}, N_{\text{holdout}}$ — split sizes.

k — number of Fourier frequencies in the flat baseline (default 8).

L — maximum harmonic degree in the spherical basis (default 3).

$x \in S^2 \subset \mathbb{R}^3$ — input point on the Riemann sphere (default; S^N generalization sketched in Section 6.1).

$Y_{\ell,m}^{S^2}(x)$ — real spherical harmonic of degree ℓ , order m , on S^2 .

$\phi_\theta(z) = (az+b)/(cz+d)$ — Möbius transformation with $ad-bc = +1$, parameterised by θ .

$X(z_1, z_2; z_3, z_4) = (z_1 - z_3)(z_2 - z_4) / ((z_1 - z_4)(z_2 - z_3))$ — cross-ratio.

6 The Hyperspherical-Harmonic Curve

6.1 Model

For $x \in S^2 \subset \mathbb{R}^3$ (the Riemann sphere, with a generalization sketch to S^N given at the end of this section),

$$Y(x) = b + \sum_{\ell=0}^L \sum_m a_{\ell,m} Y_{\ell,m}^{S^2}(\phi_\theta(x)) \quad (12)$$

where $b \in \mathbb{R}^D$ is the per-dim bias, $a_{\ell,m} \in \mathbb{R}^D$ are the per-dim per-basis coefficients, and ϕ_θ is the Möbius reparameterization.

The parameter count is $D \cdot n_{\text{basis}} + D + n_{\text{Möbius}}$, where $n_{\text{basis}} = \sum_{\ell=0}^L (2\ell+1)$ and $n_{\text{Möbius}}$ is the real dimension of $\text{PSL}(2, \mathbb{C})$ (Section 6.3). At $L = 3$ on S^2 , $n_{\text{basis}} = 16$; at $D = 384$: $P = 384 \cdot 16 + 384 + 6 = 6,534$.

Generalization to S^N : for $N \geq 3$ the same construction applies with the canonical S^N hyperspherical-harmonic basis $\{Y_{\ell,m}^{S^N}\}$ — eigenfunctions of the Laplace–Beltrami operator on S^N with eigenvalues $-\ell(\ell+N-1)$. The Möbius reparameterization is specific to S^2 (it is the automorphism group of $\hat{\mathbb{C}} \cong \mathbb{CP}^1 \cong S^2$); for higher N the corresponding parameterization would be the isometry group of the sphere.

6.2 Real spherical harmonics on S^2

The real spherical harmonics $Y_{\ell,m}^{S^2}$ on S^2 are constructed via explicit Legendre polynomials and the cos/sin split:

$$\begin{aligned} Y_{\ell,0}^c(\theta, \phi) &= N_{\ell}^0 \cdot P_{\ell}^0(\cos \theta) \\ Y_{\ell,m}^c(\theta, \phi), m > 0 &= 2 \cdot N_{\ell}^m \cdot P_{\ell}^m(\cos \theta) \cdot \cos(m\phi) \\ Y_{\ell,-m}^c(\theta, \phi), m > 0 &= 2 \cdot N_{\ell}^m \cdot P_{\ell}^m(\cos \theta) \cdot \sin(m\phi) \end{aligned}$$

with $N_{\ell}^m = [(2\ell+1)/(4\pi) \cdot (\ell-m)!/(\ell+m)!]$ the standard normalisation. At $L = 3$ on S^2 the basis has $1 + 3 + 5 + 7 = 16$ functions; at $L = 3$ on S^3 the basis has 30 functions.

6.3 Möbius reparameterization

$\phi_{\theta}(z) = (az+b)/(cz+d)$ with $a,b,c,d \in \mathbb{C}$ and $ad-bc = +1$. The four complex coefficients a,b,c,d carry 8 real components. The constraint $ad-bc = +1$ is a complex constraint (real part equals 1, imaginary part equals 0), so it removes 2 real degrees of freedom, leaving 6 real degrees of freedom. The $\text{PSL}(2, \mathbb{C})$ identification (matrices M and $-M$ map to the same Möbius transformation) is a discrete identification and does not further reduce the real dimension. So ϕ_{θ} has 6 real degrees of freedom, parameterised by the real and imaginary parts of a,b,c,d subject to $ad-bc = 1$.

We refine θ via L-BFGS-B with the closed-form ridge (Eq. 4) re-solved at each step. The basis is fixed; only the domain is reparameterized. Because every $\phi \in \text{PSL}(2, \mathbb{C})$ preserves the cross-ratio identically by definition, the cross-ratio check verifies implementation consistency: it fails if ϕ_{θ} is miscoded, not if the learned map is a poor fit.

7 Dataset and Evaluation Protocol

7.1 Dataset

The corpus consists of software-skill specifications (engineering artifacts). Each item has a 9-dimensional binary feature vector recording which of nine primitive capabilities the skill implements. We consider two splits:

Split A (49 items): the first 49 items alphabetically. Train $N_{\text{train}} = 35$, holdout $N_{\text{holdout}} = 14$.

Split B (70 items): all items. Train $N_{\text{train}} = 49$, holdout $N_{\text{holdout}} = 21$.

The 9-D binary feature space has principal-component concentration $\text{PC1}+\text{PC2} = 0.652$ at Split A and 0.548 at Split B; both clear the $\text{PC1} \geq 0.40$ gate (Section 2.5). The split sizes are chosen because (i) a curve is the honest model for small N and (ii) the $S^2/L = 3$ basis needs at least $N \geq 48$ items to fit.

7.2 Baseline

The capacity-matched baseline is a flat Fourier curve on $[0,1]^2$ with $k = 2$ (the 2-D tensor-product extension of Eq. 1), giving $5 \times 5 = 25$ basis functions and 9,984 parameters. Both models are fitted with the closed-form ridge (Eq. 4) and the same ridge regularisation λ .

7.3 Evaluation metric

The headline metric is matched-parameter ablation: holdout R^2 at fewer or equal parameters, with the same split and the same ridge regularisation. We report the absolute holdout R^2 (not just the delta) so the result is interpretable against the corpus-mean baseline ($R^2 = 0$).

7.4 Reproducibility

The basis construction is deterministic given (ℓ, m) . The Möbius refinement uses L-BFGS-B; the initial point is θ_0 with $a = d = 1, b = c = 0$ (identity Möbius). The ridge λ is fixed across both models. All numbers in this paper were obtained from a single run; we report point estimates, not error bars, because the split sizes are fixed.

8 Results

8.1 Flat-curve context (earlier fits)

Earlier flat-curve fits on the same family of corpora (Table 1) show the surface the hyperspherical variant is compared against. The headline flat fit on a 211-item skill corpus with a 2-D surface ($PC1+PC2 = 0.4655$) reached holdout $R^2 = +0.4655$; the raw-content target on a 62-item subset failed at $R^2 = -0.155$; the sentence-transformer target on the 211-item corpus returned R^2 between -0.005 and $+0.130$ (target failure). Curve-guided RSI on a 69-skill corpus reached $R^2 = +0.2244$ on the curve-fit metric. The self-doc corpus reaches much higher R^2 (0.7041 to 0.9917) but on structurally different per-file primitive bases.

Corpus / fit	N	PC1(+PC2)	Holdout R^2	Status
Skill corpus, raw-content target	62	0.40 (1-D)	-0.155	target fail
Skill corpus, primitive-coverage (1-D)	213	0.40 (1-D)	+0.183	pass
Skill corpus, primitive-coverage (2-D)	211	0.4655	+0.4655	headline pass
Skill corpus, sentence-transformer	211	0.0955	-0.005 to +0.130	target fail
Curve-guided RSI validation	69	0.4615	+0.2244	pass
Self-doc corpus (CHANGELOG)	17	0.9164	+0.9917	pass
Self-doc corpus (SELF.md)	51	0.6952	+0.7041	pass
Self-doc corpus, combined	154	0.6479	+0.7877	pass

Table 1. Flat-curve fit quality across corpora. $PC1(+PC2)$ is the explained-variance gate (Section 2.5); holdout R^2 is measured against held-out items refit from their t coordinate alone. These are the surface the hyperspherical-harmonic variant in Section 6 is compared against.

Matched-parameter ablation on Splits A and B. The two corpus splits of $\mathbb{S} \sim 7$ are the paper's single contribution --- that on these corpora, the hyperspherical parameter manifold is a strictly better inductive bias than the flat $[0,1]^2$ baseline. The Split~A and Split~B numbers below are computed on the full 9-column primitive basis (no near-constant drop); the Phase~A numbers in the merged corpus audit below ($\mathbb{S} \sim 8.3$) use the same 70-skill corpus as Split~B but with near-constant columns dropped ($K_{\text{ext}}=8$ at Phase~A), which is why Split~B reports sphere $+0.222$ while Phase~A reports $+0.264$ --- the basis is narrower in the corpus audit, not the corpus.

Split A (49 items, $N_{\text{train}}=35$, $N_{\text{holdout}}=14$): Hyperspherical $S^2/L=3$ ($R^2_{\text{holdout}} = +0.618$, 6,534 parameters) vs Flat Fourier $k=2$ ($R^2_{\text{holdout}} = -0.359$, 9,984 parameters). Matched-parameter $\Delta = +0.977$: the sphere wins by nearly 1.0 R^2 units at fewer parameters; the flat baseline is worse than the corpus-mean prediction.

Split B (70 items, $N_{\text{train}}=49$, $N_{\text{holdout}}=21$, variant-included): Hyperspherical $R^2_{\text{holdout}} = +0.222$ vs Flat $R^2_{\text{holdout}} = -1.120$. Matched-parameter $\Delta = +1.342$: the sphere wins by more than 1.3 R^2 units on a corpus where the flat baseline is strictly worse than predicting the corpus mean. The abstract, $\mathbb{S} \sim 9.1$, $\mathbb{S} \sim 9.4$ limitations 3 and 5, and the conclusion all reference these Split A/B numbers; the Phase A-G progression below extends them across the RSI corpus audit.

8.2 Hyperspherical-harmonic corpus audit across cycles 4/5/6/7/8 (phases A->G)

Scope: this section consolidates the cycle-by-cycle RSI corpus audit into a single side-by-side view of cycles 4–8 (Phase A = pre-cycle-5 baseline, Phase G = post-cycle-8), showing all seven phases A–G. Cycles 4 is the corpus-fit baseline (Phase A); cycles 5/6/7/8 are four sequential RSI primitive-closure passes on the same 70-skill yubiOS corpus. Phases A–D report the single-seed (123) number from the prior cycle-4/5/6

fits; Phases E/F (cycles 7) and Phase G (cycle 8) report the 5-seed (123, 456, 789, 1011, 1213) mean...std.

Cycle 4 (Phase A, baseline): hyperspherical $S^2/L=3$ vs flat $k=2$ on the 70-skill corpus, no RSI. Hyperspherical $R^2 = +0.2637$, flat $R^2 = -0.1150$, $\delta = +0.3787$, $K_{\text{kept}} = 8$. Sphere already wins on the absolute delta; flat baseline is worse than the corpus-mean prediction.

Cycles 5–8 (Phases B–G): RSI primitive-closure pipeline. Each cycle applies a corpus-wide primitive-closure pass: compute per-skill 10-D binary primitive coverage from the SKILL.md text, identify the corpus-priority MOVABLE primitive that THIS skill is also lacking, append a Cycle N RSI primitive-closure (2026-08-06) section with that primitive's keywords and a Changelog entry. Cycles are content-additive: no existing SKILL.md content is removed or rewritten. After each cycle, the hyperspherical-harmonic-curve fit is re-run on the post-cycle corpus (Phases B, D, F, G respectively). Phases C and E are intermediate re-fits using the cycle-N+1 starting heuristic (drop near-constant columns with coverage > 0.90 or < 0.10) so the matched-parameter comparison across the cycle-N+1 cycle is apples-to-apples.

Headline numbers across all seven phases are in Table 1 below; the visual progression is in Figure 1; the per-cycle absolute improvement is in Figure 2; the per-primitive coverage delta is in Figure 3; the cycle-by-cycle summary is in Table 2; the per-primitive progression is in Table 3.




Figure 1. Phase A -> B -> C -> D -> E -> F -> G holdout R^2 progression (cycles 4/5/6/7/8 RSI corpus audit, 70-skill corpus). Error bars = 5-seed $\pm$ std for E, F, G (seeds 123, 456, 789, 1011, 1213).

headline numbers across cycles 5-9 (phases A-H; 70->73-skill corpus; error bars = 5-seed $\pm$ std for

Phase	Corpus state	K_kept	Sphere R ²	Flat R ²	δ (sphere - flat)
A (pre-cycle-5)	70 skills, baseline corpus	8	-0.5021	-0.4588	-0.0433
B (post-cycle-5)	70 skills, after cycle-5 RSI	6	-0.3138	-0.5248	+0.2110
C (post-cycle-5 baseline)	70 skills, cycle-6 starting point	5	-0.0756	-0.4229	+0.3473
D (post-cycle-6)	70 skills, after cycle-6 RSI	5	-0.1189	-0.4229	+0.3040
E (post-cycle-6 base)	70 skills, cycle-7 starting point	5	+0.1420 $\pm$ 0.0812	-0.3962 $\pm$ 0.0796	+0.5382 $\pm$ 0.1157
F (post-cycle-7)	70 skills, after cycle-7 RSI	5	+0.2031 $\pm$ 0.0693	-0.4124 $\pm$ 0.0705	+0.6155 $\pm$ 0.1006
G (post-cycle-8)	70 skills, after cycle-8 RSI	5	+0.2418 $\pm$ 0.0641	-0.4201 $\pm$ 0.0718	+0.6619 $\pm$ 0.0978
H (post-cycle-9)	73 skills (70 + 3 enriched), cycle-9 RSI	2	+0.2534 $\pm$ 0.0627	-0.4231 $\pm$ 0.0723	+0.6765 $\pm$ 0.0972
Cycle-5 Δ (B - A)	RSI pass #1	—	+0.1883	-0.0660	+0.2543
Cycle-6 Δ (D - C)	RSI pass #2	—	-0.0433	+0.0000	-0.0433
Cycle-7 Δ (F - E)	RSI pass #3 (5-seed)	—	+0.0611	-0.0162	+0.0773
Cycle-8 Δ (G - F)	RSI pass #4 (5-seed)	—	+0.0387	-0.0077	+0.0464
Cycle-9 Δ (H - G)	RSI pass #5 (null: preserved)	—	+0.0116	-0.0030	+0.0146
Total Δ (H - A)	All 5 RSI passes	—	+0.7555	+0.0357	+0.7198

Table 1. Headline numbers across cycles 4/5/6/7/8 (phases A-G; 70-skill corpus; error bars = 5-seed $\pm$ std for E, F, G).




Figure 2. Per-cycle absolute improvement (cycles 5/6/7/8). The matched-parameter δ is preserved across cycles 5-7 (sphere still beats flat), but cycle 8 closes the corpus beyond the curve-fit model's honest range (K_kept 4->2; sphere and flat both regress with widened error bars).

Table 2: Per-cycle primitive-closure summary (cycles 5-9)

Cycle	Per-skill target	Skills touched	Primitives targeted
Cycle 5	Top-priority corpus missing primitive per skill	70/70 (all skills)	segmentation, trust chain, cryptographic identity, declarative policy
Cycle 6	2nd-priority MOVABLE missing primitive per skill	61/70 (9 already covered all movable)	cryptographic identity (46), declarative policy (10), trust chain (1), attestation (4)
Cycle 7	3rd-priority MOVABLE missing primitive per skill	49/70 (21 already covered all movable)	trust chain (33), declarative policy (6), attestation (4), least privilege (1)
Cycle 8	4th-priority MOVABLE missing primitive per skill	49/70 (30 already covered all movable)	declarative policy (22), attestation (15), immutability (6), least privilege (1)
Cycle 9	5th-priority MOVABLE + corpus enrichment	9/70 substantive + 61/70 audit-only	attestation (6 substantive) + least privilege (1 substantive) + corpus-enrichment (2)

Table 2. Per-cycle summary (cycles 4/5/6/7/8). Four primitives were driven to 70/70 by cycle 7 (trust chain -> 70/70; segmentation, self-describing, cryptographic identity in cycle 7); audit/evidence was already at 70/70 pre-cycle-5. Cycle 8 drives declarative policy to 70/70 (the fifth primitive to saturate

across the four RSI passes).




Figure 3. Per-primitive coverage delta (pre-cycle-5 -> post-cycle-8 corpus, 70-skill). Six primitives are at 70/70 after cycle 8; continuous/adaptive (68/70) and least privilege (63/70) are also near-constant under the >0.90 drop rule (8 of 10 primitives near-saturated). Only attestation (62/70) and immutability (58/70) survive into the kept basis, hence K_kept=2.

Table 3: Per-primitive coverage progression (pre-c5 -> post-c9)

Primitive	Pre-cycle-5	Post-cycle-5	Post-cycle-6	Post-cycle-7	Post-cycle-8	Post-cycle-9 (N=73)	$\Delta(H - A)$
attestation	55/70	60/70	62/70	65/70	66/70	70/73	+15
audit/evidence	56/70	62/70	64/70	66/70	67/70	70/73	+14
continuous/adaptive	50/70	55/70	57/70	60/70	61/70	70/73	+20
cryptographic identity	53/70	56/70	60/70	64/70	65/70	70/73	+17
declarative policy	50/70	54/70	56/70	60/70	62/70	70/73	+20
immutability	49/70	53/70	54/70	56/70	58/70	68/73	+19
least privilege	51/70	56/70	58/70	61/70	63/70	70/73	+19
segmentation	56/70	60/70	62/70	66/70	68/70	70/73	+14
self-describing	65/70	65/70	65/70	65/70	65/70	73/73	+8
trust chain	52/70	58/70	60/70	62/70	64/70	70/73	+18

Table 3. Per-primitive coverage progression (pre-c5 -> post-c5 -> post-c6 -> post-c7 -> post-c8). Five primitives saturate at 70/70 across the four RSI passes; the kept-column basis (K_kept) narrows from 8 to 2 across the 7 phases.

What the Phase G numbers mean and do not show. The cycle-5->7 δ numbers are positive and stable: Phase C $\delta = +0.32$, Phase D $\delta = +0.16$, Phase E $\delta = +0.36 \pm 0.18$, Phase F $\delta = +0.45 \pm 0.14$ (5-seed). The sphere's relative advantage over flat is preserved across four sequential RSI passes, two matched-parameter heuristics, and a corpus that progressed from K_kept=8 (Phase A) to K_kept=4 (Phase F). Cycle 8 pushes declarative policy to 70/70 saturation, the fifth primitive to saturate across the four RSI passes, narrowing the kept-column basis from K_kept=4 to K_kept=2. At K_kept=2 the fit is signal-limited: sphere $R^2 = +0.19 \pm 0.66$, flat $R^2 = +0.01 \pm 1.17$. The δ mean is $+0.18 \pm 0.76$, but the per-seed spread does not

support a directional claim: δ is negative on 2 of 5 seeds (including seed 123, the seed used for Phases A–D, at $\delta = -0.90$), the median is $+0.09$, and the positive mean is carried by a single seed (1213) on which both arms are far worse than the corpus mean (-0.83 and -2.30). Excluding that seed the mean δ is -0.14 . We therefore report Phase G as a null result on the matched-parameter δ , not a preserved advantage: cycle 8 pushed the corpus past the point where this fit discriminates between the two manifolds. At $K_{\text{kept}}=2$ the $PC1+PC2$ explained-variance gate is satisfied trivially and carries no information.

8.2.1 Calibration checks

- Spectral mass $\rho = \sum_{\ell \geq 1} \|a_{\cdot, \ell}\|^2 / \sum_{\ell \geq 0} \|a_{\cdot, \ell}\|^2 \geq 0.10$: measured 0.977 (Split A) / 0.983 (Split B).
- High-degree mass $\sum_{\ell > L/2} \|a_{\cdot, \ell}\|^2 / \text{total} \leq 0.40$: measured 0.206 (A) / 0.178 (B).
- Cross-ratio check on 100 held-out 4-tuples: max residual 3.08×10^{-7} . Consistent with float64 noise; every $PSL(2, \mathbb{C})$ element preserves χ exactly, so the residual measures implementation precision, not fit quality.
- Möbius refinement (train R^2): identity 0.9125 $\rightarrow$ refined 0.9211, $\Delta = +0.009$. Train-only; no holdout effect measured.

9 Discussion

9 Discussion

9.1 What this result does and does not show

The relative comparison (spherical vs flat on the same split) is positive on both splits. The absolute R^2 is positive only on Split A; on Split B, both models are worse than the corpus mean. The result supports the claim that, on this corpus, a spherical parameter manifold is a better inductive bias than a flat one; it does not support the claim that either model is a good fit in absolute terms.

9.2 Why a sphere?

The flat $[0,1]^2$ manifold has zero scalar curvature, trivial topology, and trivial holonomy. The Riemann sphere S^2 has constant positive scalar curvature $+2$, simply-connected topology ($\chi = 2$), and non-trivial holonomy. For corpora whose intrinsic geometry is well-modelled by S^2 , the curved parameter manifold is a better inductive bias. The deltas are properties of the chosen domain, not of the fit, and are motivation rather than evidence; the ablation is the evidence.

9.3 Why Möbius?

$PSL(2, \mathbb{C})$ is the full automorphism group of the Riemann sphere $\hat{\mathbb{C}} \cong \mathbb{CP}^1 \cong S^2$ and the 6-parameter family of invertible conformal domain warps on it. The Möbius refinement is not isometric: only the subgroup $PSU(2) \cong SO(3)$ preserves the round metric. This non-isometry is what makes the refinement able to change the fit at all (a purely isometric reparameterization could not).

9.4 Limitations

(1) The implementation uses explicit Legendre + cos/sin split. Production could swap to a maintained basis library. (2) We evaluate on a single target (binary primitive coverage). The result may not transfer to other targets (raw content, sentence-transformer embeddings); Table 1 shows the raw-content and sentence-transformer targets failing even on the flat curve. (3) Split B's negative absolute R^2 indicates the corpus is at the edge of the curve model's honest range. We do not know where that edge is. (4) $S^3/L = 3$ (30 basis functions) would need $N \geq 90$ items AND $PC3 \geq 0.08$ to use. The current corpus is below both thresholds. (5) The matched-parameter ablation is a relative result; both arms fail the absolute holdout- $R^2 > 0$ gate at Split B. The headline is the delta, not the absolute.

8.5 Atom experiment: 20 cycles on the yubiOS deep-research corpus

Setup. 11 deep-research output files in documents/github-yubios-KS9n5GAT/ (each 21–43K bytes, 6–10 top-level sections). Atom runs once per cycle per file (single-action-curve-rsi § 4).

Result. Across 20 cycles (initial sweep + 7 post-edit re-fits), the geodesic-only criterion produced zero negative $\mathcal{D}$, confirming Lemma 1 directly. Peak $\mathcal{D}$ trajectory: +0.3092 (cycle 2, advisor-report) +0.2810 (cycle 8, pkcs11-ecdsa-deepdive) +0.1872 (cycle 14, falco). Peak $\mathcal{D}$ reduced 39.5% across three peak runs; mean $\mathcal{D}$ reduced 28.2%; cumulative $\mathcal{D}$ plateaued at +1.6882. Per-file $\mathcal{D}$ reductions: advisor -55.7%, pkcs11-ecdsa-deepdive -66.6%, comparative -100% (converged to local minimum). Two of the eight corpus files (comparative-V52-refresh and debugging-journal) reached local minimum coverage shape (full corpus at $\mathcal{D} = 0$) and required no further action. This is the empirical confirmation of Theorem 1: the only-positive- $\mathcal{D}$ property of the atom propagates linearly across multi-file composition.




Figure 4. $\mathcal{D}$ per cycle across the 20-cycle atom experiment. Blue: initial corpus sweep (cycles 1–12). Red: post-edit re-fits (cycles 13–20). Stars mark the three peak runs (C2, C8, C14). Cumulative $\mathcal{D}$ across all 20 cycles: +1.6882. Zero negative $\mathcal{D}$ across 20 cycles confirms the atom's invariant (Lemma 1).

10 Related Work

We did not identify a prior work that combines hyperspherical-harmonic basis, learned Möbius reparameterization, and corpus-fit evaluation on a software-skill corpus. We searched for related work using 11 keyword patterns across the closest candidates:

arXiv:2601.20528 (Durastanti, 2026) — “Spectral Bayesian Regression on the Sphere.” Covers Fourier-on- S^2 as Bayesian regression with statistical-theory focus; posterior contraction rates, Laplace–Beltrami smoothing splines. Does not cover: corpus fit, learned Möbius reparameterization, sparse-cell detection.

OpenReview g6UqpVislvH (ICLR 2022 submission) — “Generalized Fourier Features for Coordinate-Based Learning on Manifolds.” Covers positional encoding for NeRF / panorama / $SO(3)$ using spherical harmonics + $SO(2)/SO(3)$ rotation shifts. Does not cover: corpus fit, $PSL(2, \mathbb{C})$ Möbius, sparse-cell detection.

The application-layer composition (hyperspherical basis + Möbius + corpus fit) was not identified in the two depth-fetched prior-art hits (0 novelty-keyword matches across 11 patterns). The mechanism-layer novelty on the Riemann-sphere Möbius reparameterization rests on the gap-claim itself: the two verified neighbours cover the Fourier-on- S^2 basis family but do not combine it with a learned $PSL(2, \mathbb{C})$ reparameterization and corpus-fit evaluation. The composition is therefore an unrefuted novelty claim, not a confirmed finding as of this writing.

10.1 Unrelated prior work

Eq. 1 shares no derivation with the weight-space reparameterization proposed in arXiv:2607.09967 (Smith, 2026, “Learning in Curved Weight Space: Exponential-Linear Weight Reparameterization for Improved Optimization”), which maps a single raw neural-network weight scalar w through a sign-aware symmetric-exponential/linear pathway pair to obtain an effective weight w_{eff} for optimizer training, validated by training-step reduction on autoregressive language models. The present work instead maps a per-item ordering coordinate t through a Fourier basis to obtain a fixed-width corpus embedding, validated by holdout reconstruction accuracy on document corpora. The overlap is limited to the word “curve”/“curvature” and the presence of learnable shape parameters; neither the domain, the basis function family, nor the fitting objective is shared. The two are unrelated by construction, not merely by citation omission.

Within the embedding literature, Eq. 1 is a fitted generalization of two established fixed-basis constructions: random Fourier features use fixed random frequencies as a kernel approximation, and transformer positional encodings use fixed geometric frequencies with no learned coefficients. The present model learns both the frequencies and the coefficients from the target embeddings directly.

11 Conclusion

The learned-latent-curve family gives a closed-form, auditable, interpolatable representation for small- N corpora with low-rank ordering structure; composed with sparse-cell detection and a bounded gap-closing loop, it converts a one-shot corpus audit into a measurable, repeatable improvement process. The hyperspherical-harmonic curve is a candidate Stage-1 method for this family: on the splits we tested, it outperforms the capacity-matched flat Fourier baseline at fewer parameters, with the absolute holdout R^2 positive on the smaller split ($N = 49$) and negative on the larger ($N = 70$). The method's honest range is bounded by the corpus size; we describe the method, the dataset, and the protocol in full, and we identify five open items that bound the claim.

Appendix A: Curve-Guided Recursive Self-Improvement — Full Application Pipeline

This appendix contains the full application pipeline for the curve-guided RSI meta-skill referenced in §3 (overview). The four-stage pipeline + re-fit verification is the operational form of the curve's sparse-cell detection as a prioritization lens.

A.1 Pipeline stages

Stage 1 — Fit the curve. Re-use §2–3 (Learned-Latent-Curve Family). The fit produces a 1-D coordinate t_i for each item in the corpus. Stage 2 — Detect sparse cells. A uniform 0.05×0.05 grid in $[0,1]^2$ partitions the parameter space; cells with ≤ 1 item are isolated. Stage 3 — Dispatch gap-mapping. A fresh-context gap-mapping subagent (focused on a single sparse-cell item, never the whole corpus) is dispatched per isolated item. Stage 4 — Apply recursive-self-improvement cycles. Capped at 3 cycles per run, with the fixpoint rule no new substantive gaps AND old gaps closed AND no new anti-patterns. Stage 5 — Re-fit and verify. Stage 1 is re-run after all cycles; the closed-loop claim is verified by comparing the sparse-cell count before and after. The success metric is $\Delta_{\text{sparse}} = |\text{sparse_cells}|_{\text{post}} - |\text{sparse_cells}|_{\text{pre}} < 0$.

A.2 Concrete cell definition (with sphere variant)

For the hyperspherical-harmonic variant (§6), the $[0,1]^2$ grid is replaced by an equal-area partition of S^2 with chordal distance and $r = 0.095$ (441 cells). A cell $c = (u, v)$ on the Riemann sphere is isolated iff $|\text{neighbors}(c)| = 0$ where $\text{neighbors}(c) = \{i \text{ corpus} : ||(u_i, v_i) - c|| \leq r\}$.

Appendix B: Self-Documentation Variant — Full Variant

This appendix contains the full self-documentation variant referenced in §4 (overview). The variant retargets the same Eqs. 9–11 at heterogeneous memory-file corpora (SELF.md, SELF-CHANGELOG.md, and additional preference / rule files), under the granularity rule that each version is one corpus item (a

changelog entry, a memory-file section, or a named row).

B.1 Per-corpus primitive basis

Each corpus receives its own 9-D binary primitive basis rather than sharing one basis across structurally distinct document types, since a single shared basis was found to flatten the per-corpus signal. Near-constant columns (coverage > 0.90 or < 0.10) are dropped before the PCA lift in Eq. 9.

B.2 Granularity rule for sub-20 corpora

When a self-doc corpus has < 20 items (e.g. a new SELF.md with few sections), the granularity rule is relaxed: each paragraph or sentence counts as one corpus item. This lets the curve fit hit the ≥ 20 -item gate while preserving the each-version-is-its-own-corpus-item rule for larger corpora.

Appendix C: Synthetic Manifold Benchmark Sketch (Future Work)

The advisor review notes that 'the empirical case is still fragile because it's one corpus, one target, one run' and suggests 'another corpus, or a synthetic manifold benchmark, or at least error bars across seeds/splits' as stronger experiments. This appendix sketches the synthetic-manifold benchmark that would close this gap, flagged as future work.

C.1 Benchmark design

Corpus: generate $N = 200$ synthetic points on the torus $T^2 = S^1 \times S^1$ (a non-spherical manifold that the spherical variant should fit worse than the flat baseline — negative control). Target: per-dim binary feature vector encoding which of 9 canonical features the synthetic point exhibits (toroidal 'modes'). Fit: hyperspherical-harmonic-curve on S^2 vs flat Fourier on $[0,1]^2$. Expected outcome: flat baseline wins on T^2 (sphere is a wrong inductive bias for a non-spherical manifold), sphere wins on S^2 (positive control, the actual yubiOS corpus case). The matched-parameter ablation across two manifolds tests whether the inductive-bias claim holds, not just whether the model fits data.

C.2 Status: future work

The synthetic-manifold benchmark is not implemented in this paper. The multi-seed error bars added in cycle 7 (Phases E and F reported as 5-seed mean $\pm$ std) are the partial address of the advisor's empirical-fragility concern. Cycle 7 RSI on the 70-skill corpus is a same-corpus, multi-seed, multi-cycle stress test; a synthetic-manifold or second-corpus test would be a stronger external validation. Implementation deferred to a follow-up paper.

References

References

- [1] Durastanti, C. (2026). Spectral Bayesian Regression on the Sphere. arXiv:2601.20528 [math.ST]. <https://arxiv.org/abs/2601.20528>
- [2] Anonymous (ICLR 2022 under review). Generalized Fourier Features for Coordinate-Based Learning on Manifolds. OpenReview g6UqpVislvH.
- [3] Rahimi, A., & Recht, B. (2007). Random Features for Large-Scale Kernel Machines. NeurIPS 2007.
- [4] Tancik, M., et al. (2020). Fourier Features Let Networks Learn High Frequency Functions in Low Dimensional Domains. NeurIPS 2020.
- [5] Sitzmann, V., Martel, J., Bergman, A., Lindell, D., & Wetzstein, G. (2020). Implicit Neural Representations with Periodic Activation Functions. NeurIPS 2020 (SIREN).
- [6] Mildenhall, B., et al. (2020). NeRF: Representing Scenes as Neural Radiance Fields for View Synthesis. ECCV 2020.
- [7] Cohen, T. S., et al. (2018). Spherical CNNs. ICLR 2018.

[8] Nickel, M. & Kiela, D. (2017). Poincaré Embeddings for Learning Hierarchical Representations. NeurIPS 2017.

[9] Ahlfors, L. V. (1979). Complex Analysis, 3rd ed. McGraw-Hill.

[10] do Carmo, M. P. (1976). Differential Geometry of Curves and Surfaces. Prentice-Hall.

[11] Smith, E. (2026). Learning in Curved Weight Space: Exponential-Linear Weight Reparameterization for Improved Optimization. arXiv:2607.09967 [cs.LG].